

PLATE 2.7 · COHESION – MAJORITY ACCEPTANCE

Majority acceptance on Q21

The weighting gate, the three parked constitutional questions, and what scoring this gauge would do to Quality of Life. Nothing implemented, no config touched — this is the memo, the ruling is yours.

GATE · RESOLVE FIRST

Are the Online Analysis percentages weighted?

FINDING Yes — weighted. The gate clears, and the rest of this memo is live.

This is not a documentation answer. WVS documents nothing about it: there is no weight toggle in the tool, no weight variable anywhere in its 27-item “Cross by” list or in the full Wave 7 question index, no mention in the FAQ, and nothing in the site's own client-side JavaScript (`js/jds.js` and the inline blocks contain no weight logic at all). The prior session's finding on that point still holds exactly.

It is instead a **proof from the tool's own arithmetic**, and it is decisive in form rather than suggestive. A table of raw integer frequencies has one property it cannot violate: cells can never sum to *more* than their own marginal. Missing data pushes a sum below the marginal; nothing pushes it above. The tool's cross-tabulations break that rule repeatedly.

CELLS SUMMING ABOVE THEIR OWN MARGINAL – IMPOSSIBLE FOR RAW COUNTS

TABLE	CELL GROUP	CELLS SUM TO	STATED MARGINAL
AUS	Q21 × income, “Fifth step” column	374	373
AUS	Q21 × income, “Nineth step” column	28	27
AUS	Q21 × income, “Mentioned” row	162	161
CAN	Q21 × age, “16-24” column	437	436
CAN	Q21 × age, “25-34” column	660	659
CAN	Q21 × age, “45-54” column	721	720
CAN	Q21 × age, “Mentioned” row	353	352
USA	Q21 univariate, all four categories	2,597	2,596
USA	Q21 × income, “Don’t know” row	95	94

Every figure above was read off the tool's own “Show Counts” view and re-added by hand, independent of the parser written for the bulk sweep.

The only way a categorical frequency table acquires non-integer underlying quantities is weighting. The displayed “Number of cases” is a **rounded weighted count**, and the percentage beside it is computed on the same weighted basis.

The control that makes it airtight

If this were a parsing artefact or a server bug, it would appear everywhere. It doesn't. **Germany, Japan and South Korea return perfectly consistent tables on every crossing tried** — 12-category income, sex, three-band education, six-band age — with zero mismatches across every row, column and marginal. Australia, Canada and the United States break on some crossings and not others, which is exactly what rounding predicts and exactly what a systematic bug would not do.

That split is itself the confirming signal: WVS supplies no weight (or a weight identically 1) for some Wave 7 samples, and a real one for others. Countries in the first group can't produce a rounding artefact because there is nothing to round.

Two residuals, stated rather than papered over.

Which weight is not determinable from outside. What can be pinned is that the weighted N matches each country's true sample size to the unit (Germany 1,528; Australia 1,813; Canada 4,018; USA 2,596), so it is a nationally-equilibrated weight in the S017 family — not S018, which rescales every country to a common N and would visibly change these totals.

The microdata cross-check you listed second could not be completed, and the reason is worth recording. Every download on worldvaluessurvey.org — microdata, codebooks, the per-country technical reports, even the published “WVS Results By Country 2017-2022” PDF — sits behind a registration form requiring a real name, institution, e-mail and a signed conditions-of-use agreement. The Online Analysis tool is the *only* ungated route on the entire site. I did not submit your identity to a third party or accept a licence on your behalf; that is your call to make, not a step to take silently mid-task.

Is hand-entering this table computing, or republishing?

Something changed while testing the tool that reframes this question, so take it before the arguments: the tool's Display control offers **Show Column % (excluding DK/NA)**, and it works on a single-country table. **WVS's own tool will emit the valid-base percentage and the valid N directly.** The figure the Scorecard would enter is one the provider computed and displayed as such — not one derived here from a count and a denominator. That removes the sharpest edge of this question before it is argued.

THE CASE FOR REPUBLISHING

The provider computed the statistic. WVS's server applied the weight, chose the denominator, and rendered a percentage to one decimal. Nothing was derived here; a number was transcribed.

The transport is a red herring. **housing-pressure** is built on an SDMX key this project reverse-engineered from conflict errors; **life-satisfaction** runs on an undocumented internal API found by reading a JS bundle, complete with a signing salt. Neither was treated as a sourcing problem, because how the bytes arrived says nothing about who did the arithmetic.

And a human click-path genuinely exists: **WVSONline.jsp** is a normal page that embeds this exact tool. The POST chain is the machine route to the same screen, not a back door to something the public can't reach.

THE CASE FOR COMPUTING

Every other manual gauge transcribes from a document that *exists* — a Gallup article's table, an OECD explorer export, a PISA release. This table has no existence until a query is constructed. Country, wave, item and display basis are all inputs chosen here. Change the display basis and the published number changes.

That is a real distinction, and it is not about transport. It is about whether the provider published *this figure* or published a facility for generating figures. A one-decimal percentage is not sitting in a document anywhere; it is rendered on request.

The **.xls** export weakens this side rather than strengthening it: the file is generated from the same query, so it inherits every choice the on-screen table made. Export and screen are the same act.

RECOMMENDATION

This is republishing, and the test that draws the line should be written once, now. Proposed wording for METHODOLOGY.md:

A figure is *republished* when the provider's own system performs every statistical operation behind it — weighting, aggregation, the choice of denominator — and emits the finished figure. Selecting which of the provider's figures to read is not an operation. A figure is *computed* when this project performs any statistical step the provider did not: combining series, deriving a denominator, splicing units, or averaging items into an index the provider does not publish.

Q21 read off the tool's valid-base view passes cleanly. It also draws the line somewhere useful: it is precisely the test that **fails** a self-assembled Q121-Q130 composite (Question 3), fails combining Great Britain with Northern Ireland into a UK figure, and would have failed the Productivity unit-splice this project already refused. The rule earns its place by ruling things out, not just in.

Gallup-shaped construct, or a WVS-scoped rename?

Gallup's Migrant Acceptance Index averages three items at increasing personal proximity — migrants living in the country, becoming neighbours, marrying into the family — into a 0-9 scale where higher is more accepting. WVS Q21 is one rejection-framed item at a single proximity: *“Could you please mention any that you would not like to have as neighbors? ... Immigrants/foreign workers.”*

KEEP THE NAME, ACCEPT A PARTIAL PROXY

Q21 is the middle rung of Gallup's own ladder, so this is a genuine subset rather than a different subject. The gauge's job in Quality of Life — does the majority receive newcomers well — is served by any honest measure of it.

Renaming costs continuity: the gauge, its plate, its prose and the whole recorded search history are indexed under one name, and a second name for a third instrument makes the record harder to follow.

RENAME TO WHAT WVS MEASURES

Evidence found in this session says the proxy is not merely partial — it is **differently ordered**. Germany is the *most* accepting peer on Q21 at 3.9%, and among the *most* negative on Q121, where 31.2% call the impact of immigrants bad. The Netherlands runs the same way.

A measure that reverses a country's standing against a neighbouring item is not quietly narrower than the three-item construct. Naming it after Gallup's index would assert an equivalence this session has direct evidence against.

RECOMMENDATION — AGREEING WITH YOUR LEAN, ON STRONGER GROUNDS THAN “CALL IT WHAT IT IS”

Rename. Suggested: **“Cohesion — acceptance of migrant neighbours”**, unit % **who would not want immigrants/foreign workers as neighbours (WVS Q21)**. Keep the gauge id `cohesion-majority-acceptance` so plate, URL and history survive; rename the display name, `oneLiner` and `unit`.

The Germany/Q121 inversion is the argument. It is not a stylistic preference for precise naming — it is a concrete case where the Gallup name would make a claim about Australia's peers that the underlying item does not support.

This is a change to what the gauge measures, so CLAUDE.md's standing three-file rule binds: `gauges.config.json`, `content/why-this-matters/cohesion-majority-acceptance.md`, and its baseline in `content/why-this-matters-verification.ts` must move in one commit or the build fails.

Polarity, direction, and whether to build a composite

Polarity is `lower_is_better`, and it needs no special handling

Q21 is rejection-framed, so a lower percentage is the better outcome. The peer-relative direction basis is unaffected, and this is worth stating plainly because it looks like it should be a problem and isn't: direction

classifies the trend in Australia's *level score*, and `computeLevelScore` already applies polarity when producing that score (1 – raw for `lower_is_better`). By the time direction is computed, the number is already oriented so that up means better. The site's existing `lower_is_better` gauges — `housing-pressure`, `personal-safety`, `internal-cohesion` — all rely on exactly this and need no per-gauge handling.

Direction will read `insufficient-history`, permanently for now

Wave 7 gives each peer exactly one observation. `MIN_WAVES_FOR_TREND` is 3 and `MIN_SPAN_YEARS_FOR_TREND` is 6, so the gauge lands in the `insufficient-history` state that was built for this gauge in the first place and has never had a live trigger.

Wave 6 cannot rescue it. Checked live rather than assumed: the 2010-2014 wave carries 60 studies and **7 of our 9 peers** — Canada and Great Britain are both absent. A two-wave series would reintroduce precisely the structural incompleteness that made the Gallup data unscorable. Wave 7 is the only peer-complete wave that exists.

Q121-Q130: `exclude`, and not narrowly

The battery is peer-complete — Q121 was pulled for all nine and returns a full five-point table everywhere. Exclusion is not a coverage judgment.

It fails on **construct** and on **the test proposed in Question 1**. Read the battery's actual items: fills useful jobs, strengthens cultural diversity, increases the crime rate, gives asylum to refugees, increases terrorism risk, helps poor people, increases unemployment, leads to social conflict, plus a policy-preference item at Q130. That is perceived *consequences* of immigration and a policy stance — not whether the majority will personally accept a newcomer. It is a third construct, not a richer version of the second.

And averaging ten items into an index WVS does not publish is computing under any reading of the Question 1 test. It would mean choosing weights, choosing how to fold the reverse-coded items, and choosing what to do with the policy item — three methodology decisions, made inside a gauge, to produce a number no provider stands behind.

RECOMMENDATION

One clean item. A single provider-computed statistic the site can point at is more defensible than a composite this project assembles, and the gap between the two is exactly the compute/republish line drawn above. Q121-Q130 is worth recording as a *rejected* option with its reason, so the next session doesn't rediscover the battery and read its silence as an oversight.

The nine figures, and what they would need

WVS WAVE 7, Q21 – “WOULD NOT LIKE IMMIGRANTS/FOREIGN WORKERS AS NEIGHBOURS”

COUNTRY	FIELDWORK	% VALID BASE	VALID N	% ALL RESP.	LEVEL SCORE
DEU	2018	3.9	1,523	3.9	100.0
GBR	2022	4.7	2,585	4.7	96.9
NZL	2020	4.7	1,057	4.7	96.9
USA	2017	8.5	2,436	8.0	82.4
CAN	2020	8.8	4,018	8.8	81.3
AUS	2018	8.9	1,813	8.9	80.9
NLD	2022	20.9	2,057	20.0	35.1
KOR	2018	22.0	1,245	22.0	30.9
JPN	2019	30.1	1,311	29.1	0.0

Australia ranks 6th of 9. Level scores are min-max on the valid base with `lower_is_better`, computed with this site's own `computeLevelScore` arithmetic. The two bases differ only where a country has meaningful DK/NA — the USA (6.2%), Japan (3.1%) and the Netherlands (4.1%); four peers have none at all, which is itself the reason to prefer the valid base.

USE THE VALID BASE, AND THE REASON IS COMPARABILITY NOT PREFERENCE

The default “% of Total” view puts DK/NA in the denominator, and DK/NA rates run from 0% to 6.2% across these nine. That systematically flatters countries whose fieldwork produced more non-response. The valid-base view is provider-computed, so it costs nothing under the Question 1 test, and it removes a real cross-country artefact.

The 2017-2022 spread

Five years, USA 2017 to Netherlands and Great Britain 2022. **Acceptable, and already governed.** `scoringBasis: "latest-wave-per-country"` was built for exactly this and has sat unused since 2026-08-11; each country contributing its own most recent value is not a workaround here, it is the mechanism's original design case. The disclosure apparatus already exists: the amber alternate-basis note on the gauge page, the `/status` tag, the per-dot `asOfYear` on the dot strip and its caption line.

Worth noting the spread is genuinely narrower than the basis was built to tolerate — it was designed against Gallup waves running 2016, 2017 and 2019 with countries missing entirely.

staleAfterMonths — and the finding you should see first

NOTE Switching to WVS makes this gauge's data *older*, not fresher.

`latestDataYear` reads the last point in *Australia's own* series. Gallup's last Australian wave was 2019; Australia's WVS Wave 7 fieldwork was 2018. The gauge would move from data-through-2019 to data-through-

2018 — from roughly 80 months old to roughly 92. What the switch buys is peer-completeness (9 of 9 against 4 of 9), not currency, and the gauge copy must say so or it will read as an upgrade in both respects.

On the number itself: the gauge is already past due at its current 24 months and would remain so at any defensible threshold, since no threshold below ~92 months avoids the flag. The honest derivation available, using this project's own “real cadence, verified before use” rule, is wave-based: Australia's WVS fieldwork ran 2012 (Wave 6) → 2018 (Wave 7), a 72-month gap, plus Wave 7's roughly 24-month fieldwork-to-first-release lag, plus a small buffer — about 102 months.

My recommendation is to leave it at 24 and let it read STALE. A 102-month threshold would be arithmetic constructed to stop a true statement from displaying. The data is eight years old; the existing `staleDisclosure`, with its `{{DATA_AGE_YEARS}}` and `{{DATA_THROUGH_YEAR}}` tokens, already exists to explain that the age is structural rather than neglect. That is the disclosure this project built for exactly this shape of problem.

GB, not UK — disclosable, and quantified

A disclosure, not a disqualifier, and the case is stronger than mere consistency with the ISSP ruling. Northern Ireland is fielded as a separate Wave 7 study, so it can be measured rather than guessed at: **NIR 2022, 9.5%, N = 446** against Great Britain's 4.7%. Population-weighted, a UK-wide figure would land near 4.8%.

That difference **cannot move Australia's score at all**, because Great Britain is neither the minimum nor the maximum of the peer set — Germany and Japan are the bounds, and only the bounds and Australia's own value enter `computeLevelScore`. The exposure is confined to Great Britain's own rank, where at worst it slips from equal-second to third.

Combining the two studies ourselves would be computing under the Question 1 test. Use Great Britain alone, disclose the exclusion, and state the magnitude — which is now a measured figure rather than an unknown.

The collection procedure, and where it belongs

The framing in your brief needs one correction, and it changes the answer: **there is an ordinary UI click-path.** `WVSOnline.jsp` is a normal, human-navigable page that embeds this exact tool in an iframe. A person clicks the wave, ticks the country, picks Q21, sets the Display control, and reads or exports the table. What deep-links don't survive is URL parameters, because the shell is a JS app — so the human path is “click through from the wave selector”, not “open this URL”.

The POST chain is the *machine* route to that same screen, and the manual lane already assumes a session does the mechanical part: step 4 of `docs/manual-lane-checklist.md` reads “Hand the CSV to Claude Code.” So the fragility question resolves cleanly — nobody is being handed a POST chain.

RECOMMENDATION

Split it the way the lane is already split. The full replayable request chain goes in `data/manual/README.md`, which is the “how” file, with the three non-obvious details recorded because each one silently wasted a session: **SAIDS/AMIDS** are the split halves of the checkbox's `country:study` value; Wave 7 is **WAVE=1562**; the valid base is **CRUCES_OPERACION=2** with a trivial crossing variable. `docs/manual-lane-checklist.md` — which is explicitly schedule-only and meant to read in two minutes — gets one row naming the human click-path and pointing at `README.md`.


Re-verified live this session, not carried over: the whole chain still works on a fresh anonymous session with no login anywhere. Germany 2018 reproduced the recorded table exactly (60 / 3.9%, 1,463 / 95.7%). The .xls export returned a genuine 32 KB Composite Document File with **Content-Disposition: attachment**.

One durable warning worth recording: the WVS study IDs are not stable identifiers to hardcode. Germany appears twice in the Wave 7 list — “Germany 2018” (276:3310) and “Germany-EVS” (276:3381) — and so do the Netherlands and Great Britain. Picking the EVS twin silently returns a different survey. The study list must be re-read and matched on its **title** attribute each time, the same “discover, don’t hardcode” rule the SIPRI and OECD fetchers already follow.

IMPACT · YOU ASKED TO KNOW BEFORE DECIDING

It moves Australia's headline band

Scoring this gauge takes Quality of Life from **67.9 to 69.5**. The Strengthening/Leading boundary is 69. Australia’s headline verdict flips from **Strengthening** to **Leading**, and its rank moves from 4th to 3rd of 9.



The mechanism is simple and not a quirk: Australia scores 80.9 on this gauge against a current composite of 67.9, so an eighth gauge at that level pulls the mean up 1.6 points. The composite crosses 69 as soon as the new gauge scores above 77.

QUALITY OF LIFE COMPOSITE, AUSTRALIA – WITH AND WITHOUT THE Q21 GAUGE

MIN-MAX BOUNDS	WITHOUT BAND		WITH BAND		RANK
As published	67.9	Strengthening	69.5	Leading	4th → 3rd
USA excluded (§3.3 test)	62.8	Strengthening	65.0	Strengthening	5th → 3rd
JPN excluded	69.5	Leading	69.8	Leading	4th → 3rd

Exclusions applied consistently across all eight gauges, per §3.3’s own method — not to the new gauge alone.

My read: not a reason for caution about the data, and a real reason for caution about the copy

The band flip should carry no weight in your ruling. Three things say so, and the third is the one that settles it.

1. It does not survive the site’s own robustness test. Under §3.3’s USA-excluded bounds the composite reads Strengthening with or without the gauge. Under JPN-excluded bounds it reads Leading either way. The flip exists only at the published bounds.

2. The gauge's own band is robust — 80.9 as published, 72.4 with Japan out of the bounds, Leading under both. So this is not a **bandRobustness: "overstates"** case in the sense **life-expectancy** and **personal-safety** are. The sensitivity is in the composite's proximity to a boundary, not in the gauge.
3. The site already tells readers this. The live homepage carries a proximity disclosure reading: *"Australia, New Zealand, Netherlands and Japan all sit within a typical year's movement of the Strengthening/Leading boundary — on this measure the composite does not meaningfully separate them."* A 1.6-point move across that boundary is exactly the movement the page has already declared meaningless. Treating it as a reason to withhold a gauge would be acting against a disclosure this project published on purpose.

Withholding a real, peer-complete measure because of where it leaves the headline would be the exact failure the maturity apparatus, the exclusion disclosure and the unscored-gauge mechanism were all built to prevent. The reason to score this gauge is that it closes a genuine gap in Quality of Life; the reason not to would have to be about the data, and after the weighting finding there isn't one.

The copy is where real caution belongs. Australia's move to Leading is an artefact of adding a gauge, not of Australia changing — and it arrives alongside a *rank* change from 4th to 3rd that holds under all three bounds treatments and is the more robust of the two effects. Both need saying plainly at the moment they land, or the next reader will reasonably conclude the country improved. Two consequences follow: the peer median also moves, 63.7 → 61.9, so any prose quoting it goes stale in the same commit; and the REGISTER draft line for this gauge and the dimension will both need re-verification against the guard.

One implementation risk to check rather than assume. An **insufficient-history** gauge has no direction, so it counts toward none of the homepage's improving / flat / deteriorating tallies — currently 3 · 2 · 2 for seven scored gauges. With eight scored gauges those counts would sum to seven. If any assertion or copy treats the three as exhaustive of the scored set, it will need adjusting. Worth a grep before building, not a blocker.

What I did, and what needs you

RESOLVED	The weighting gate. Percentages are weighted; the gate clears on evidence, not assumption.
RECOMMENDED	Score it, on a renamed WVS-scoped construct, single item, valid base, latest-wave-per-country, Great Britain disclosed, stale AfterMonths left at 24.
NEEDS YOU	All three constitutional questions. Each recommendation above is a proposal, not a decision — and the compute/republish test in particular becomes durable policy the moment it is adopted.
NOT TOUCHED	<code>gauges.config.json</code> , all gauge copy, all scoring code, <code>data/processed/</code> . Nothing implemented, per your instruction.
CHANGED	One line in <code>HANDOVER.md</code> — the <code>/about</code> Contact item, logged as intentionally suppressed with a working guard, per your first instruction.
DECLINED	Submitting your name, institution and e-mail to WVS's registration form to reach the microdata. That is an outward-facing act and yours to authorise.

If you rule yes, the shortest honest path

- 1 Rename the gauge and update the three paired files in one commit, per CLAUDE.md's standing rule.
- 2 Enter the nine valid-base figures with each country's fieldwork year; log the pull in `collection-log.csv` with the real collector.
- 3 Set `scoringBasis`: "`latest-wave-per-country`", move the gauge from `unscoredDimensions` to a `weights` entry, and renormalise Quality of Life from 1/7 to 1/8 across all eight.
- 4 Record the compute/republish test in `METHODOLOGY.md` as its own dated ruling — it outlives this gauge.
- 5 Rewrite the band and median copy for the flip, and re-run the REGISTER draft-line guard.
- 6 Retire the WVS-eligibility trigger and re-date the remaining three; add Q121-Q130 as a recorded rejection with its reason.

THE AUSTRALIA SCORECARD · CONSTITUTIONAL MEMO · COHESION-MAJORITY-ACCEPTANCE
 EVIDENCE GATHERED LIVE 2026-08-26/27 AGAINST WORLDVALUESSURVEY.ORG · NO LOGIN, NO
 PAYWALL, NO LICENCE ACCEPTED
 COMPOSITE FIGURES REPRODUCED AGAINST THE LIVE BUILD (67.9 / STRENGTHENING / 4TH OF 9)